



# **Advanced Programming in the UNIX Environment**

Process Limits and Identifiers



apue\$

## getrlimit(2) / setrlimit(2)

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlp);
```

```
int setrlimit(int resource, const struct rlimit *rlp);
```

Returns: 0 on success; -1 on error

Changing resource limits follows these rules:

- a process may change its *soft limit* to a value less than or equal to its *hard limit*
- any process can lower its *hard limit* greater than or equal to its *soft limit*
- only superuser can raise *hard limits*
- changes are per process only (which is why ulimit must be a shell built-in)

apue\$

## Process Identifiers

```
#include <unistd.h>
pid_t getpid(void);
pid_t getppid(void);
```

Process ID's are guaranteed to be unique and identify a particular executing process with a non-negative integer.

Certain processes have fixed, special identifiers. They are:

- *swapper, sched, idle* or *system*, process ID 0 – responsible for scheduling
- *init*, process ID 1 – bootstraps a Unix system, owns orphaned processes
- *pagedaemon*, process ID 2 – responsible for the VM system (some Unix systems)



apue\$

## Process Limits and Identifiers

---

Certain aspects of a process's execution are restricted via resource limits.

A resource limit is specified as a *soft* limit and a *hard* limit; only the superuser may raise a hard limit.

Resource limits are enforced *per process*.

A process further has (at least) a process ID (PID) and a parent process ID (PPID). More on these process relationships in our next videos.



# **Advanced Programming in the UNIX Environment**

## **Process Control**

```
while (getinput(buf, sizeof(buf))) {  
    buf[strlen(buf) - 1] = '\0';  
  
    if((pid=fork()) == -1) {  
        fprintf(stderr, "shell: can't fork: %s\n",  
                strerror(errno));  
        continue;  
    } else if (pid == 0) { /* child */  
        execlp(buf, buf, (char *)0);  
        fprintf(stderr, "shell: couldn't exec %s: %s\n", buf,  
                strerror(errno));  
        exit(EX_UNAVAILABLE);  
    }  
  
    /* parent waits */  
    if ((pid=waitpid(pid, &status, 0)) < 0) {  
        fprintf(stderr, "shell: waitpid error: %s\n",  
                strerror(errno));  
    }  
}  
  
exit(EX_OK);
```

fork(2)

```
#include <unistd.h>
```

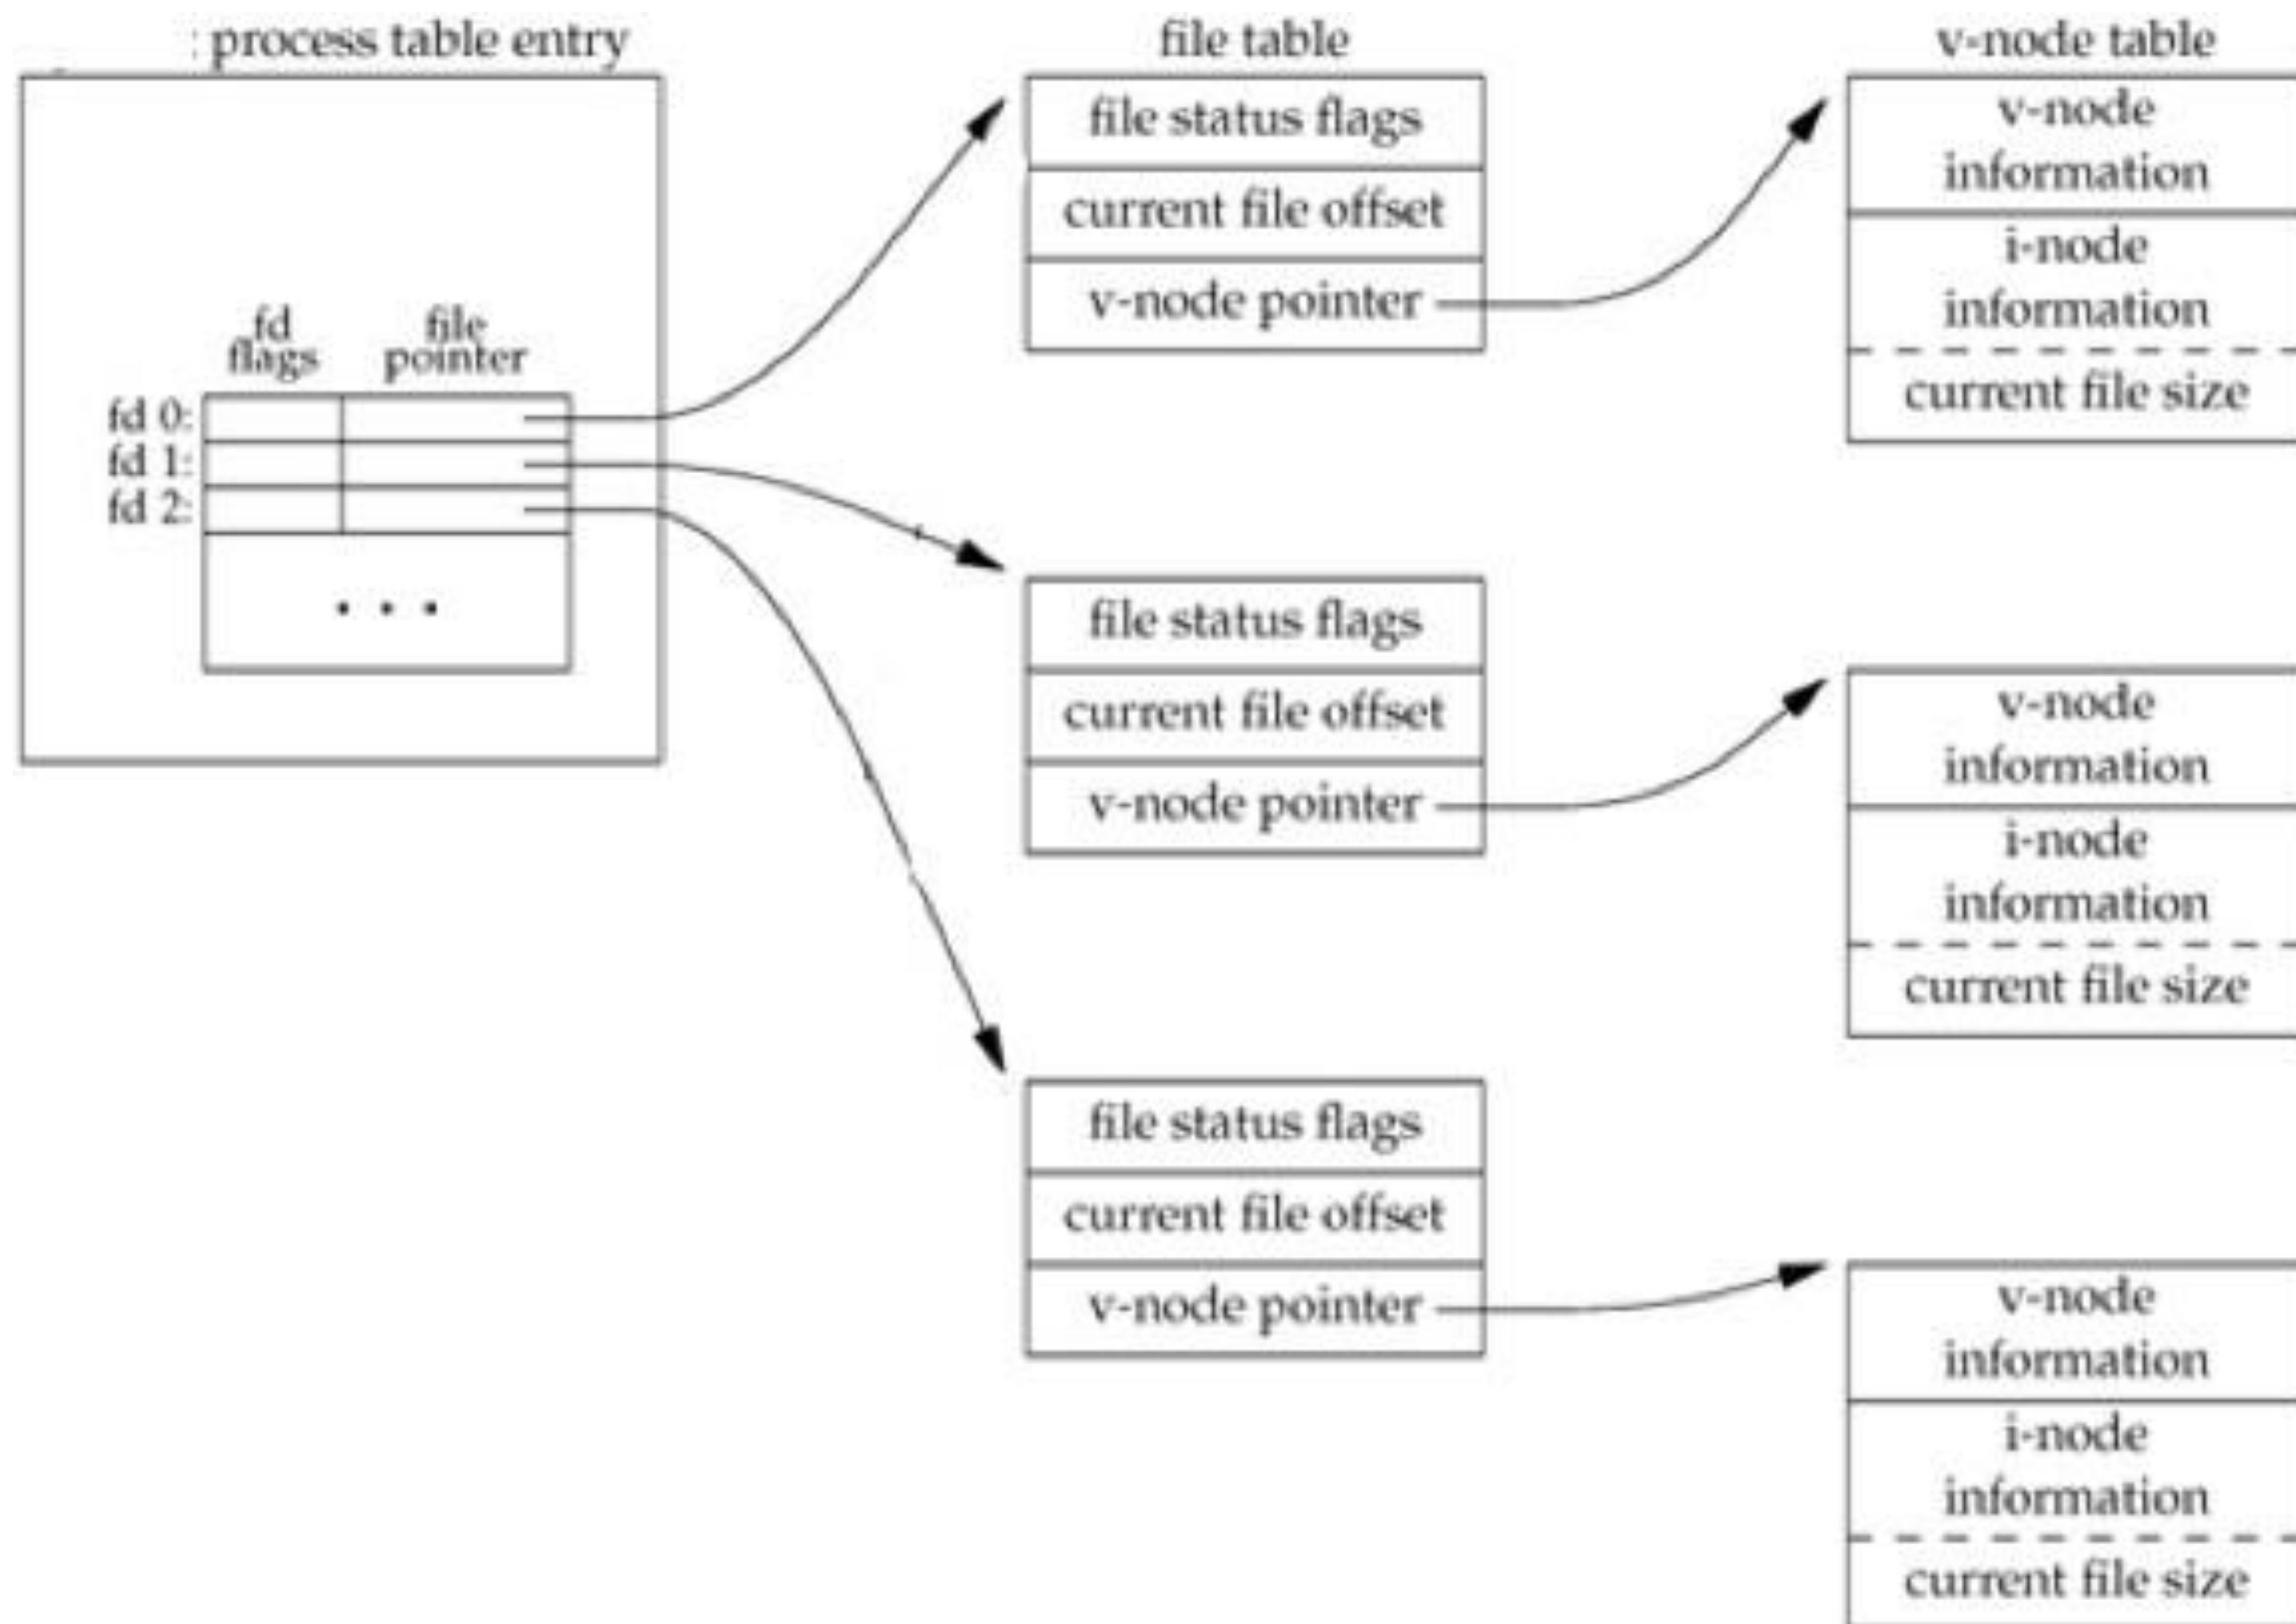
```
pid_t fork(void);
```

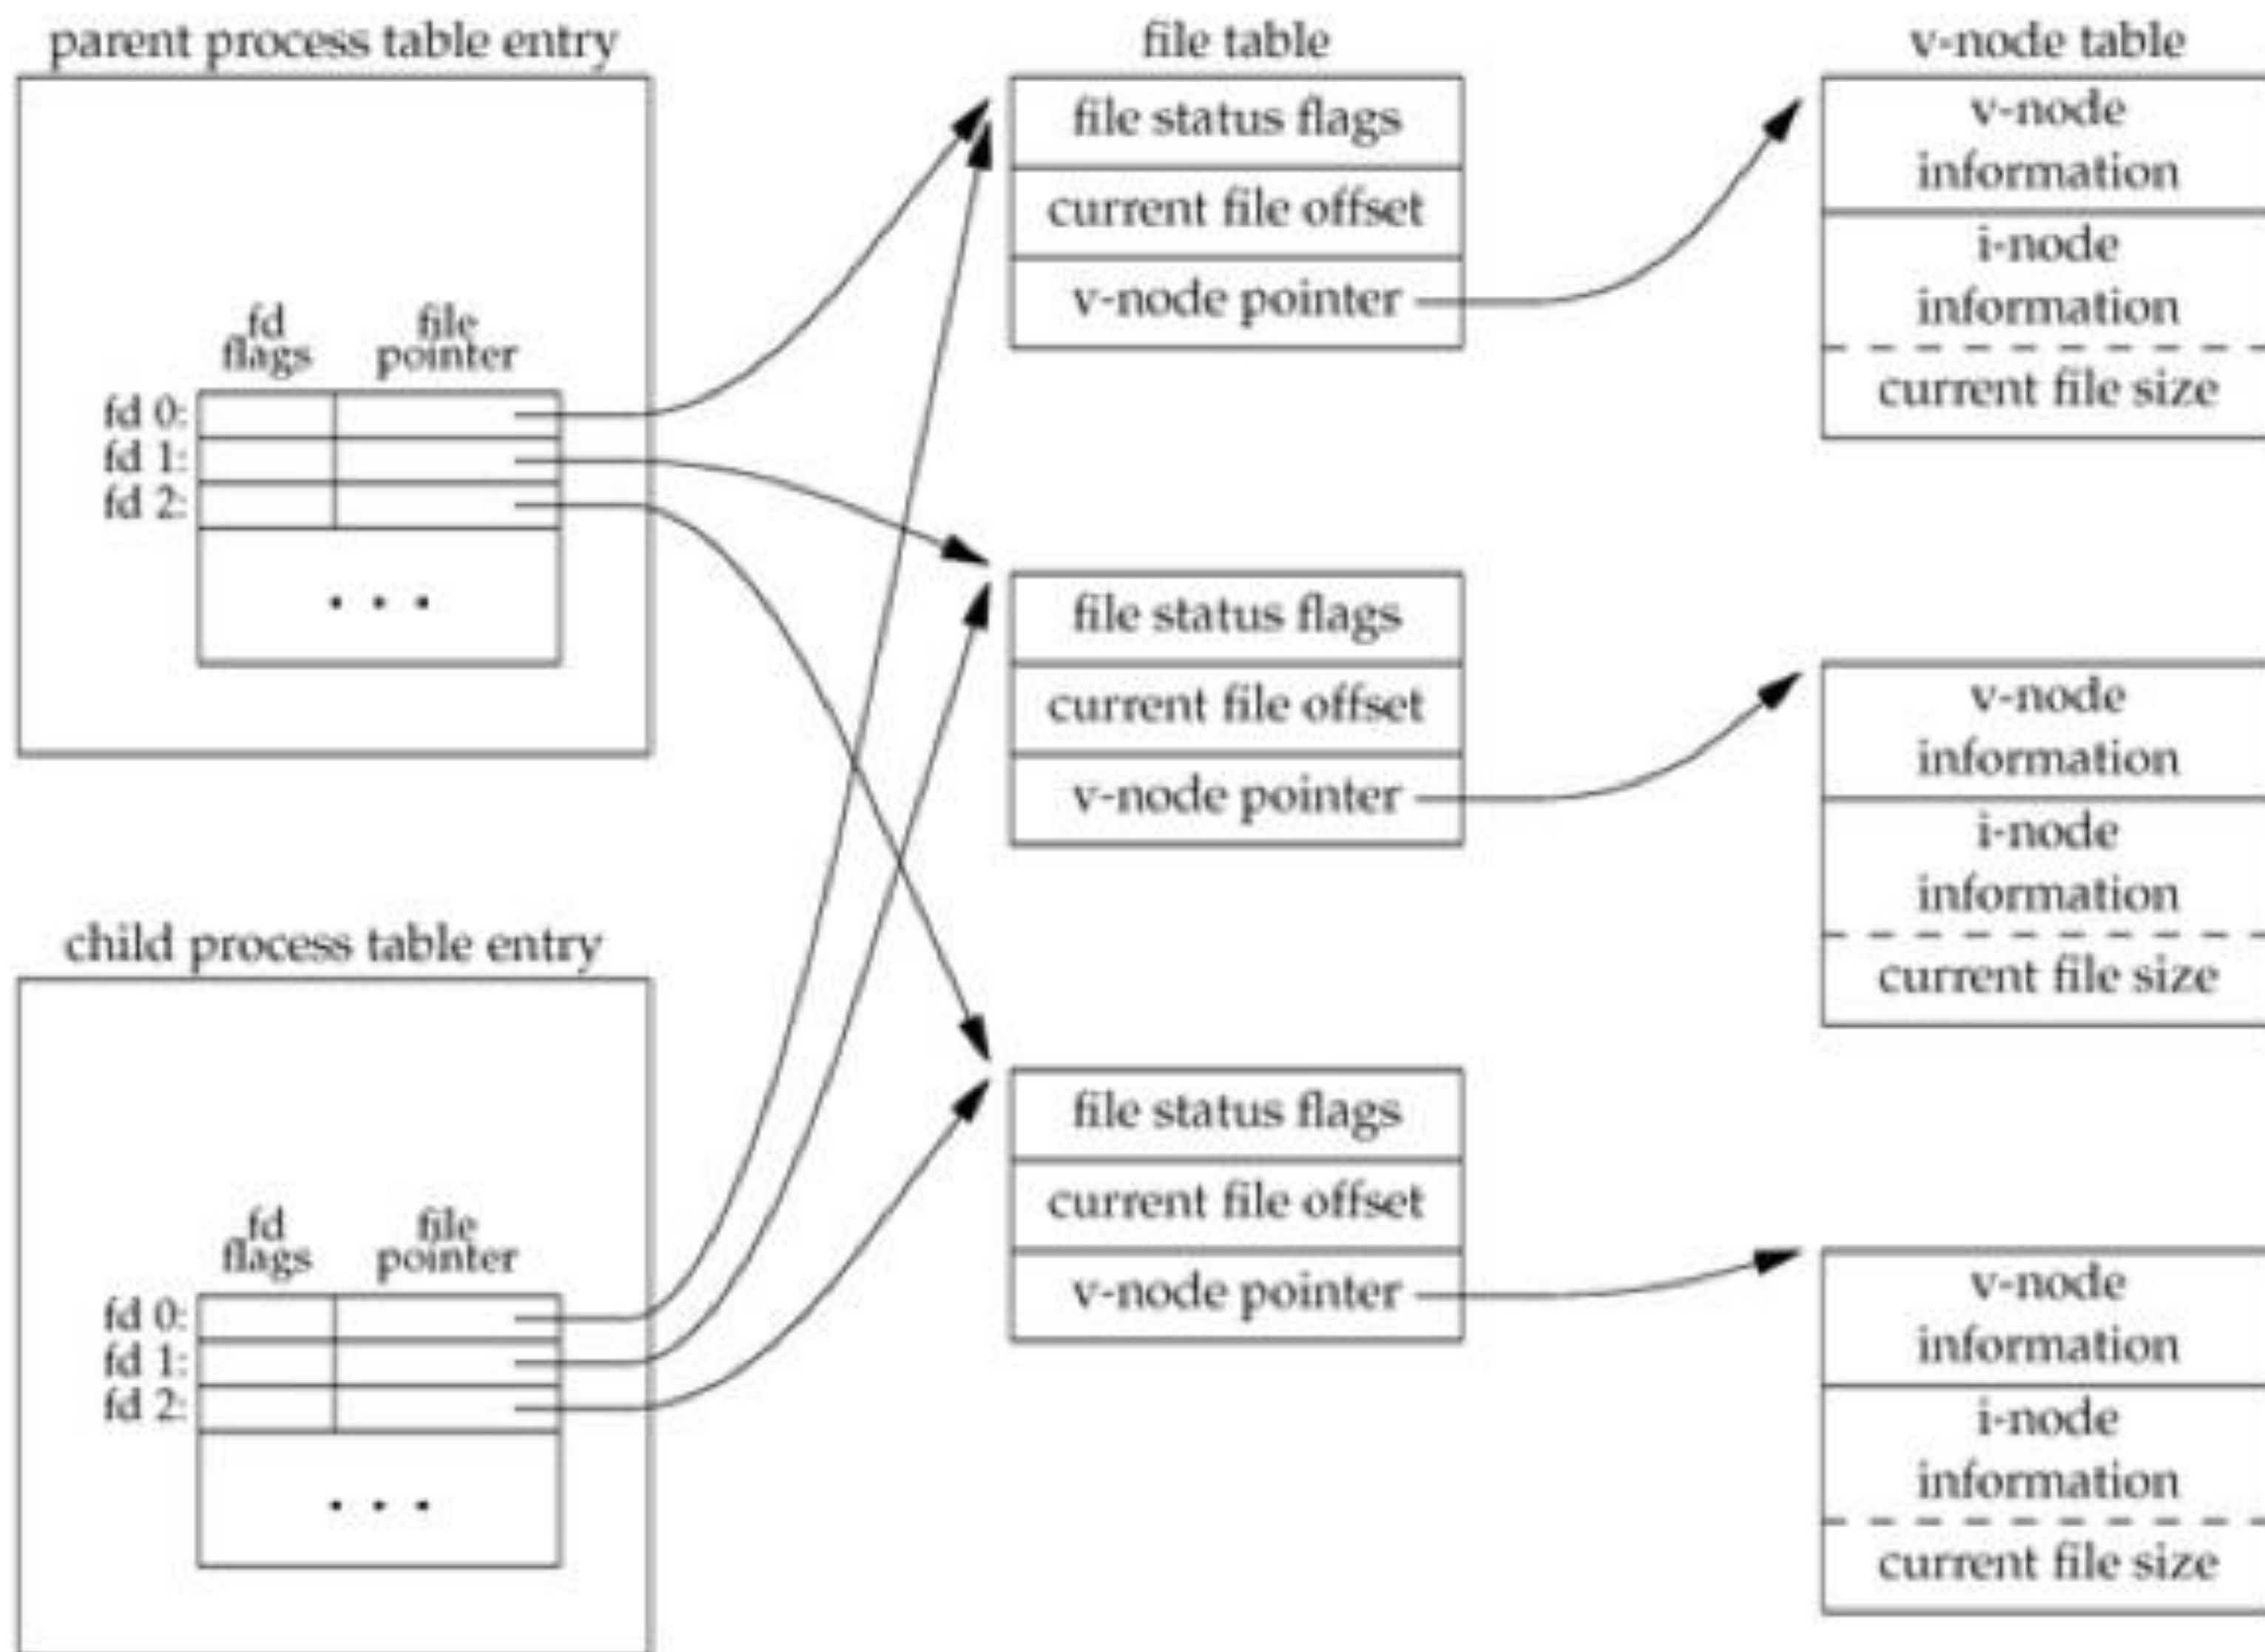
Returns: twice(!): 0 to the child, new pid to the parent; -1 on error

fork(2) causes creation of a new process. The new process (child) is an exact copy of the calling process (parent) except for the following:

- The child process has a unique process ID.
- The child process has a different parent process ID (*i.e.*, the processID of the parent process).
- The child process has its own copy of the parent's descriptors.
- The child process's resource utilizations are set to 0.

Note: no order of execution between child and parent is guaranteed!


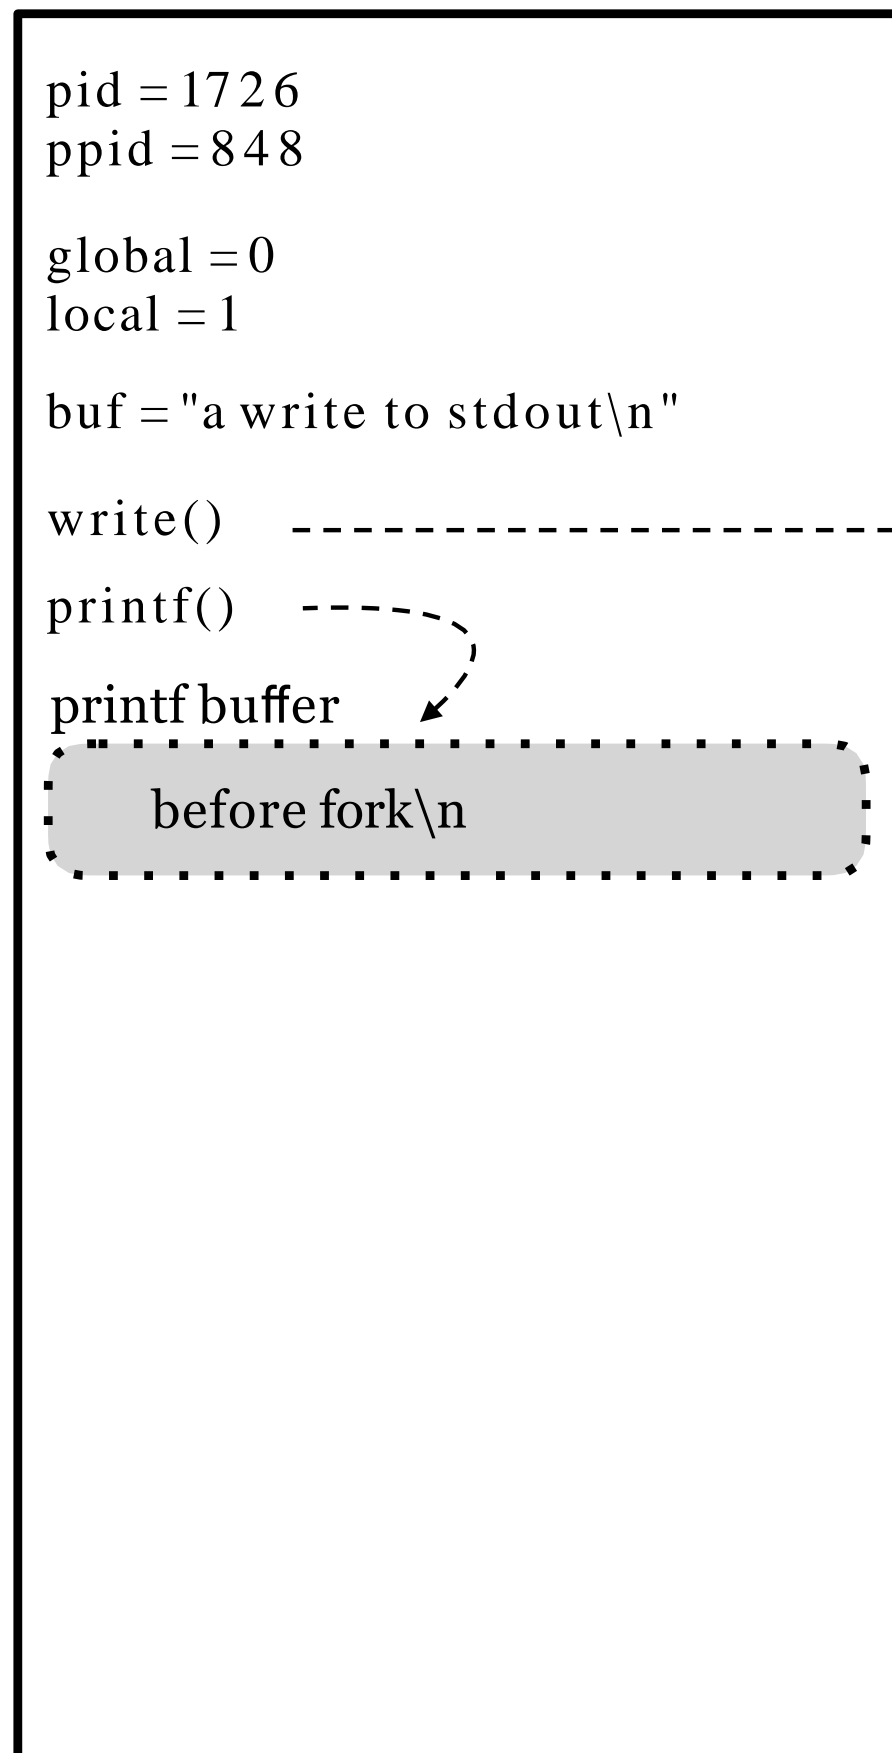





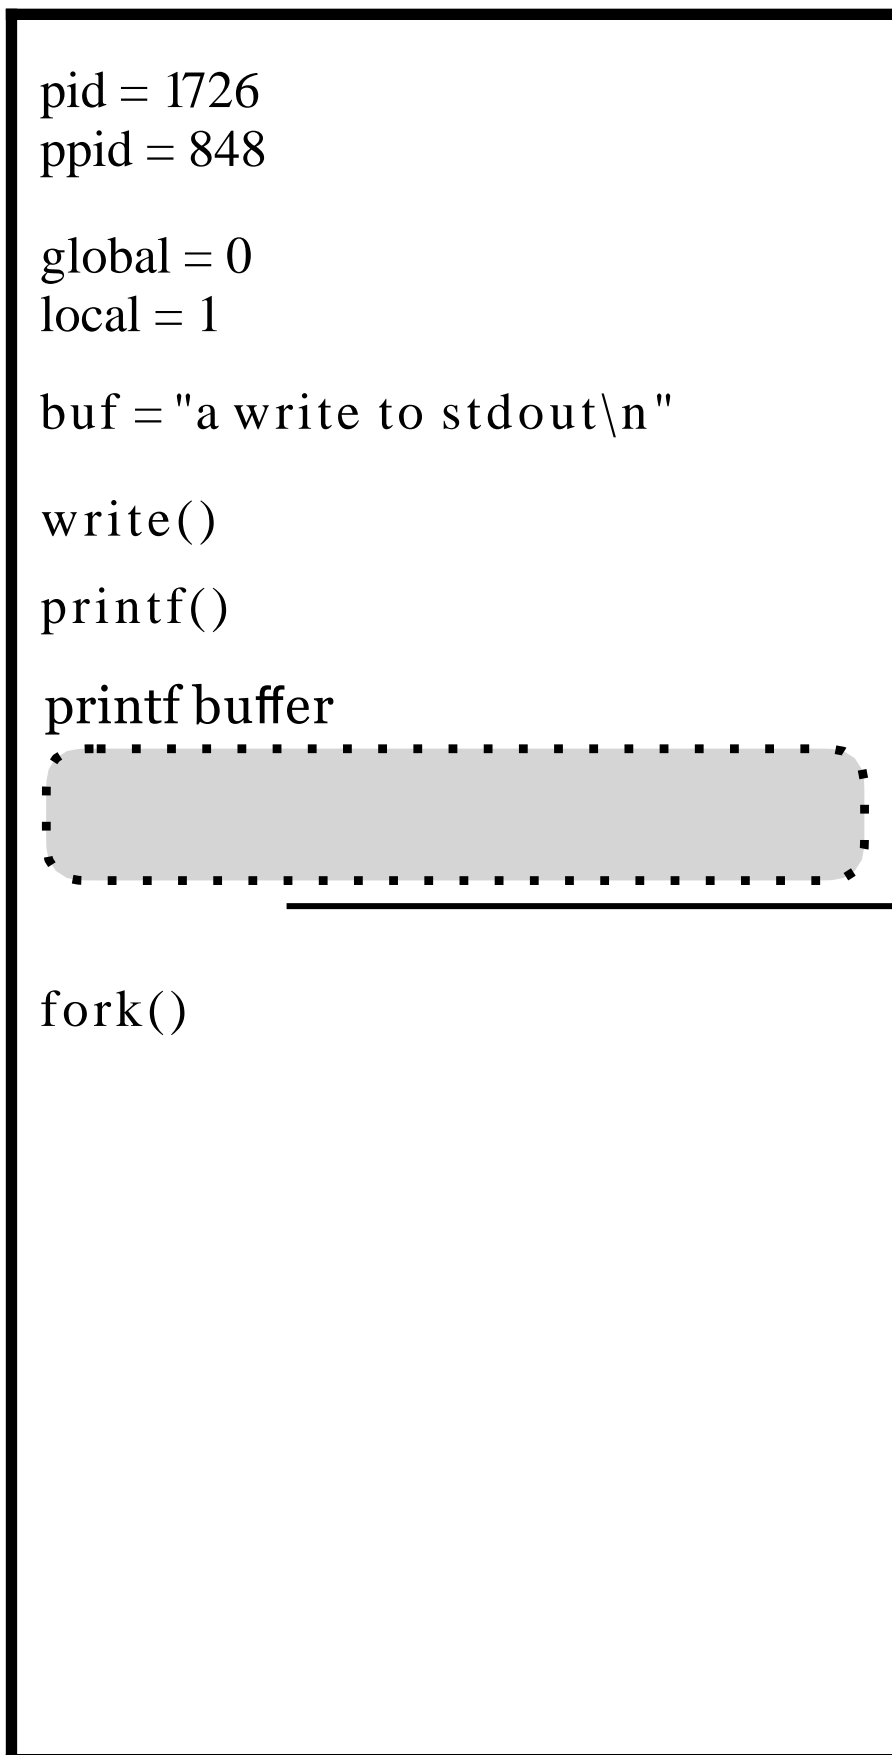


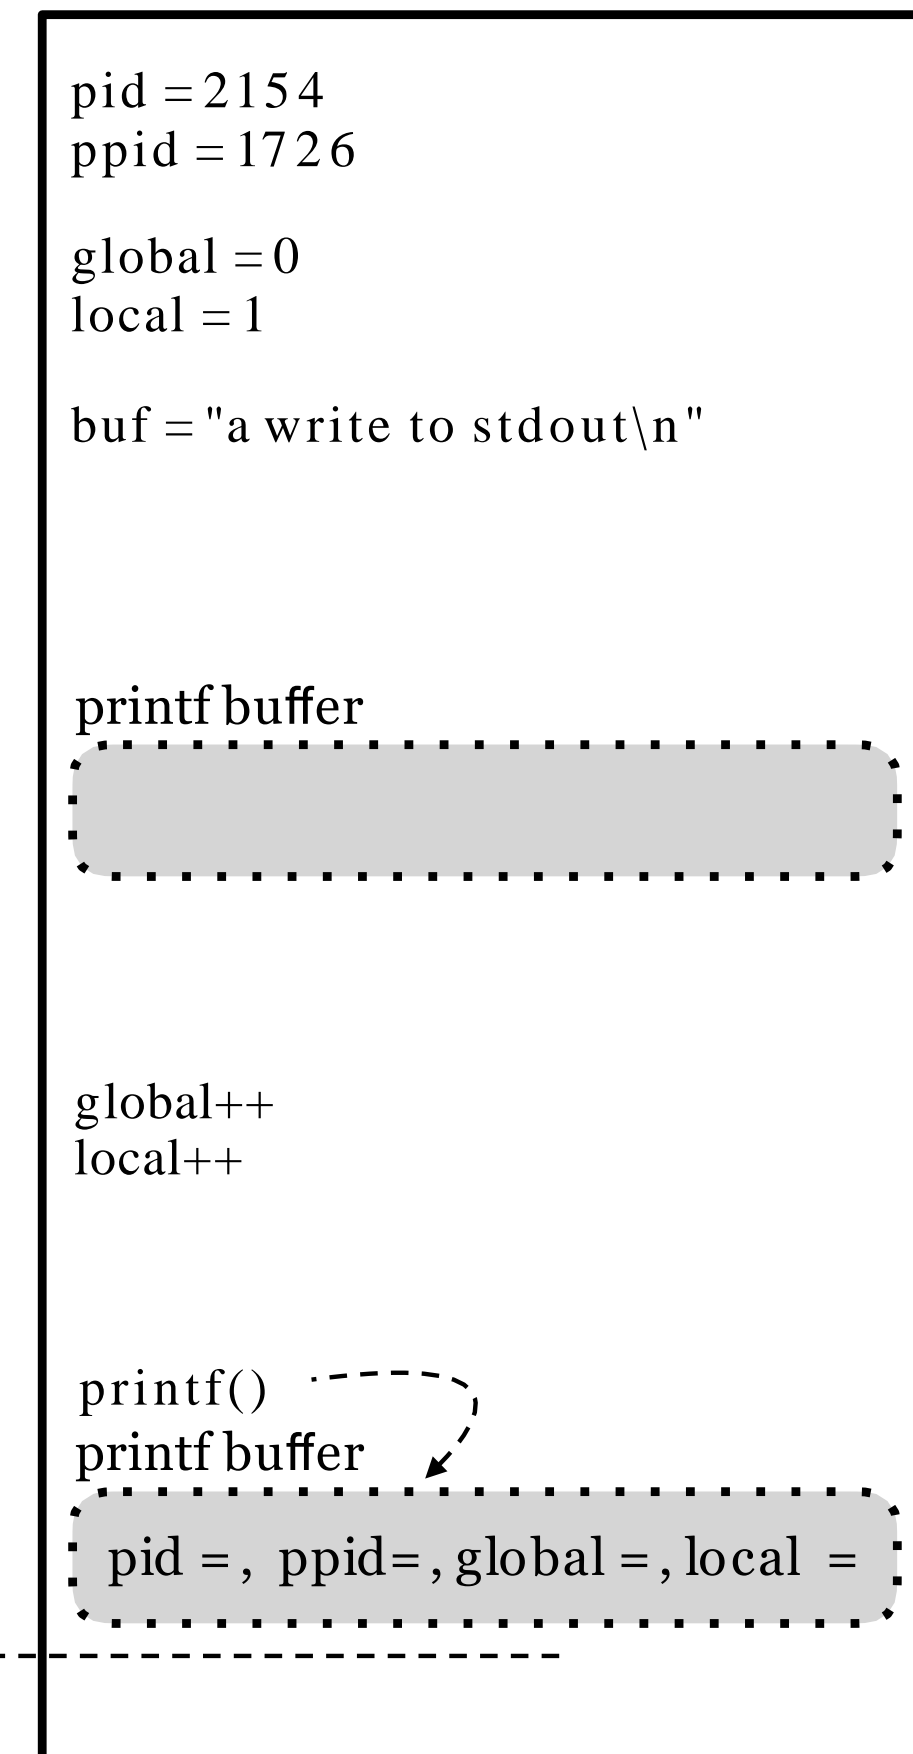

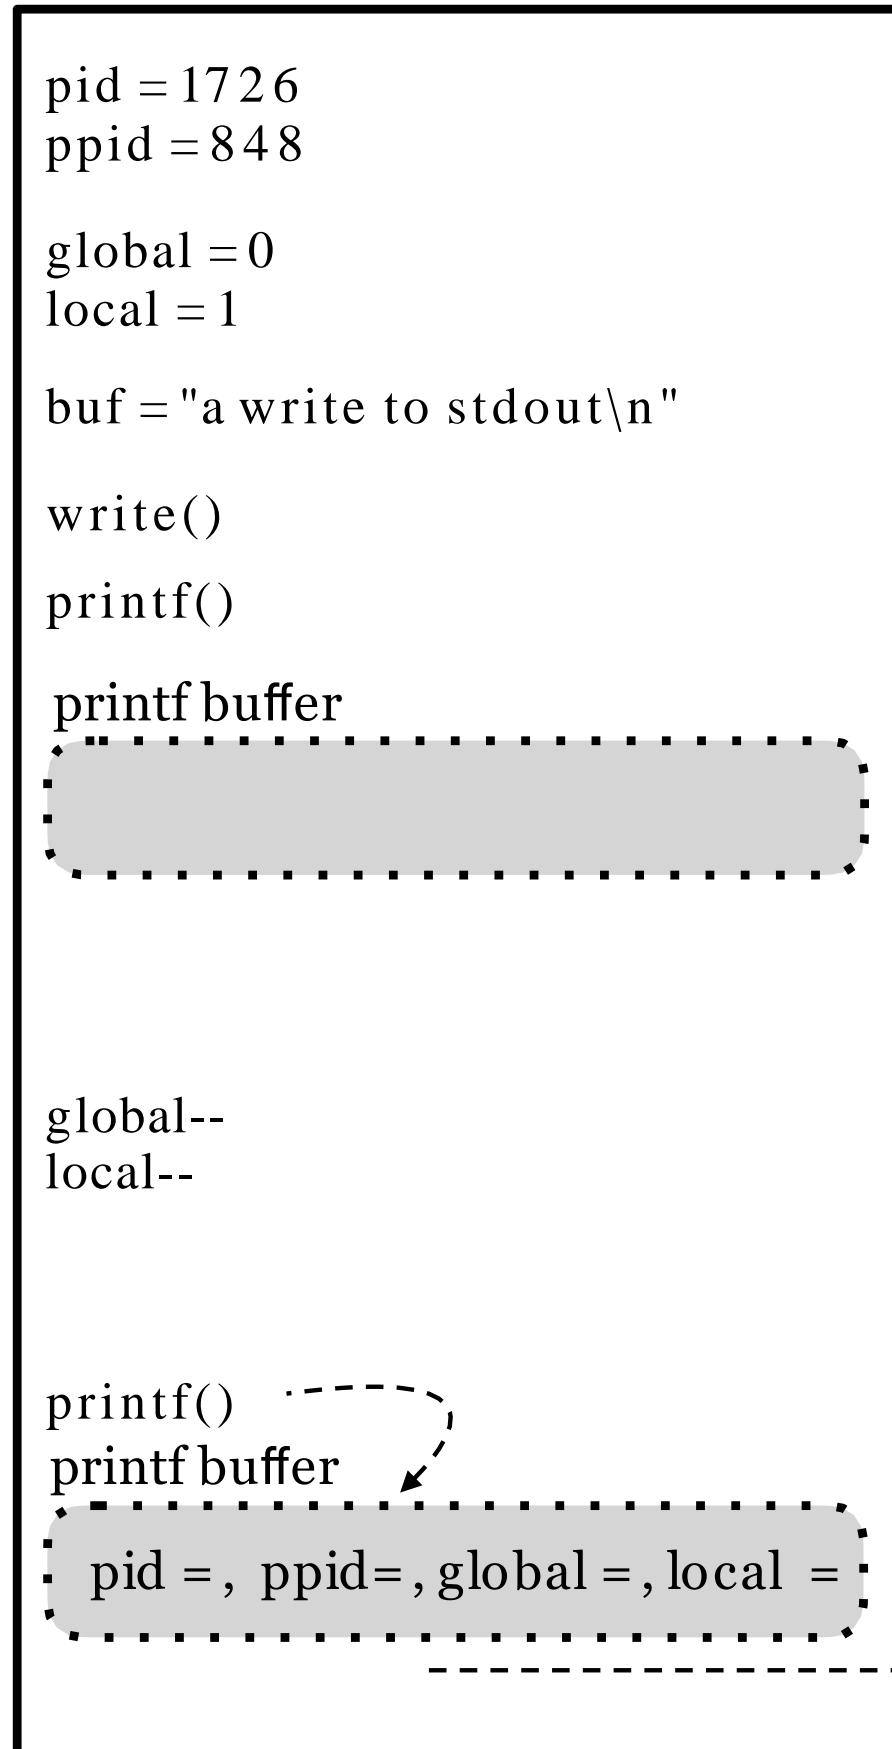
```
[apue$ vim forkseek.c  
[apue$ cc -Wall -Werror -Wextra forkseek.c  
[apue$ ./a.out forkseek.c  
Starting pid is: 361  
361 offset is now: 0  
child 999 done seeking  
361 offset is now: 64  
999 offset is now: 96  
apue$
```

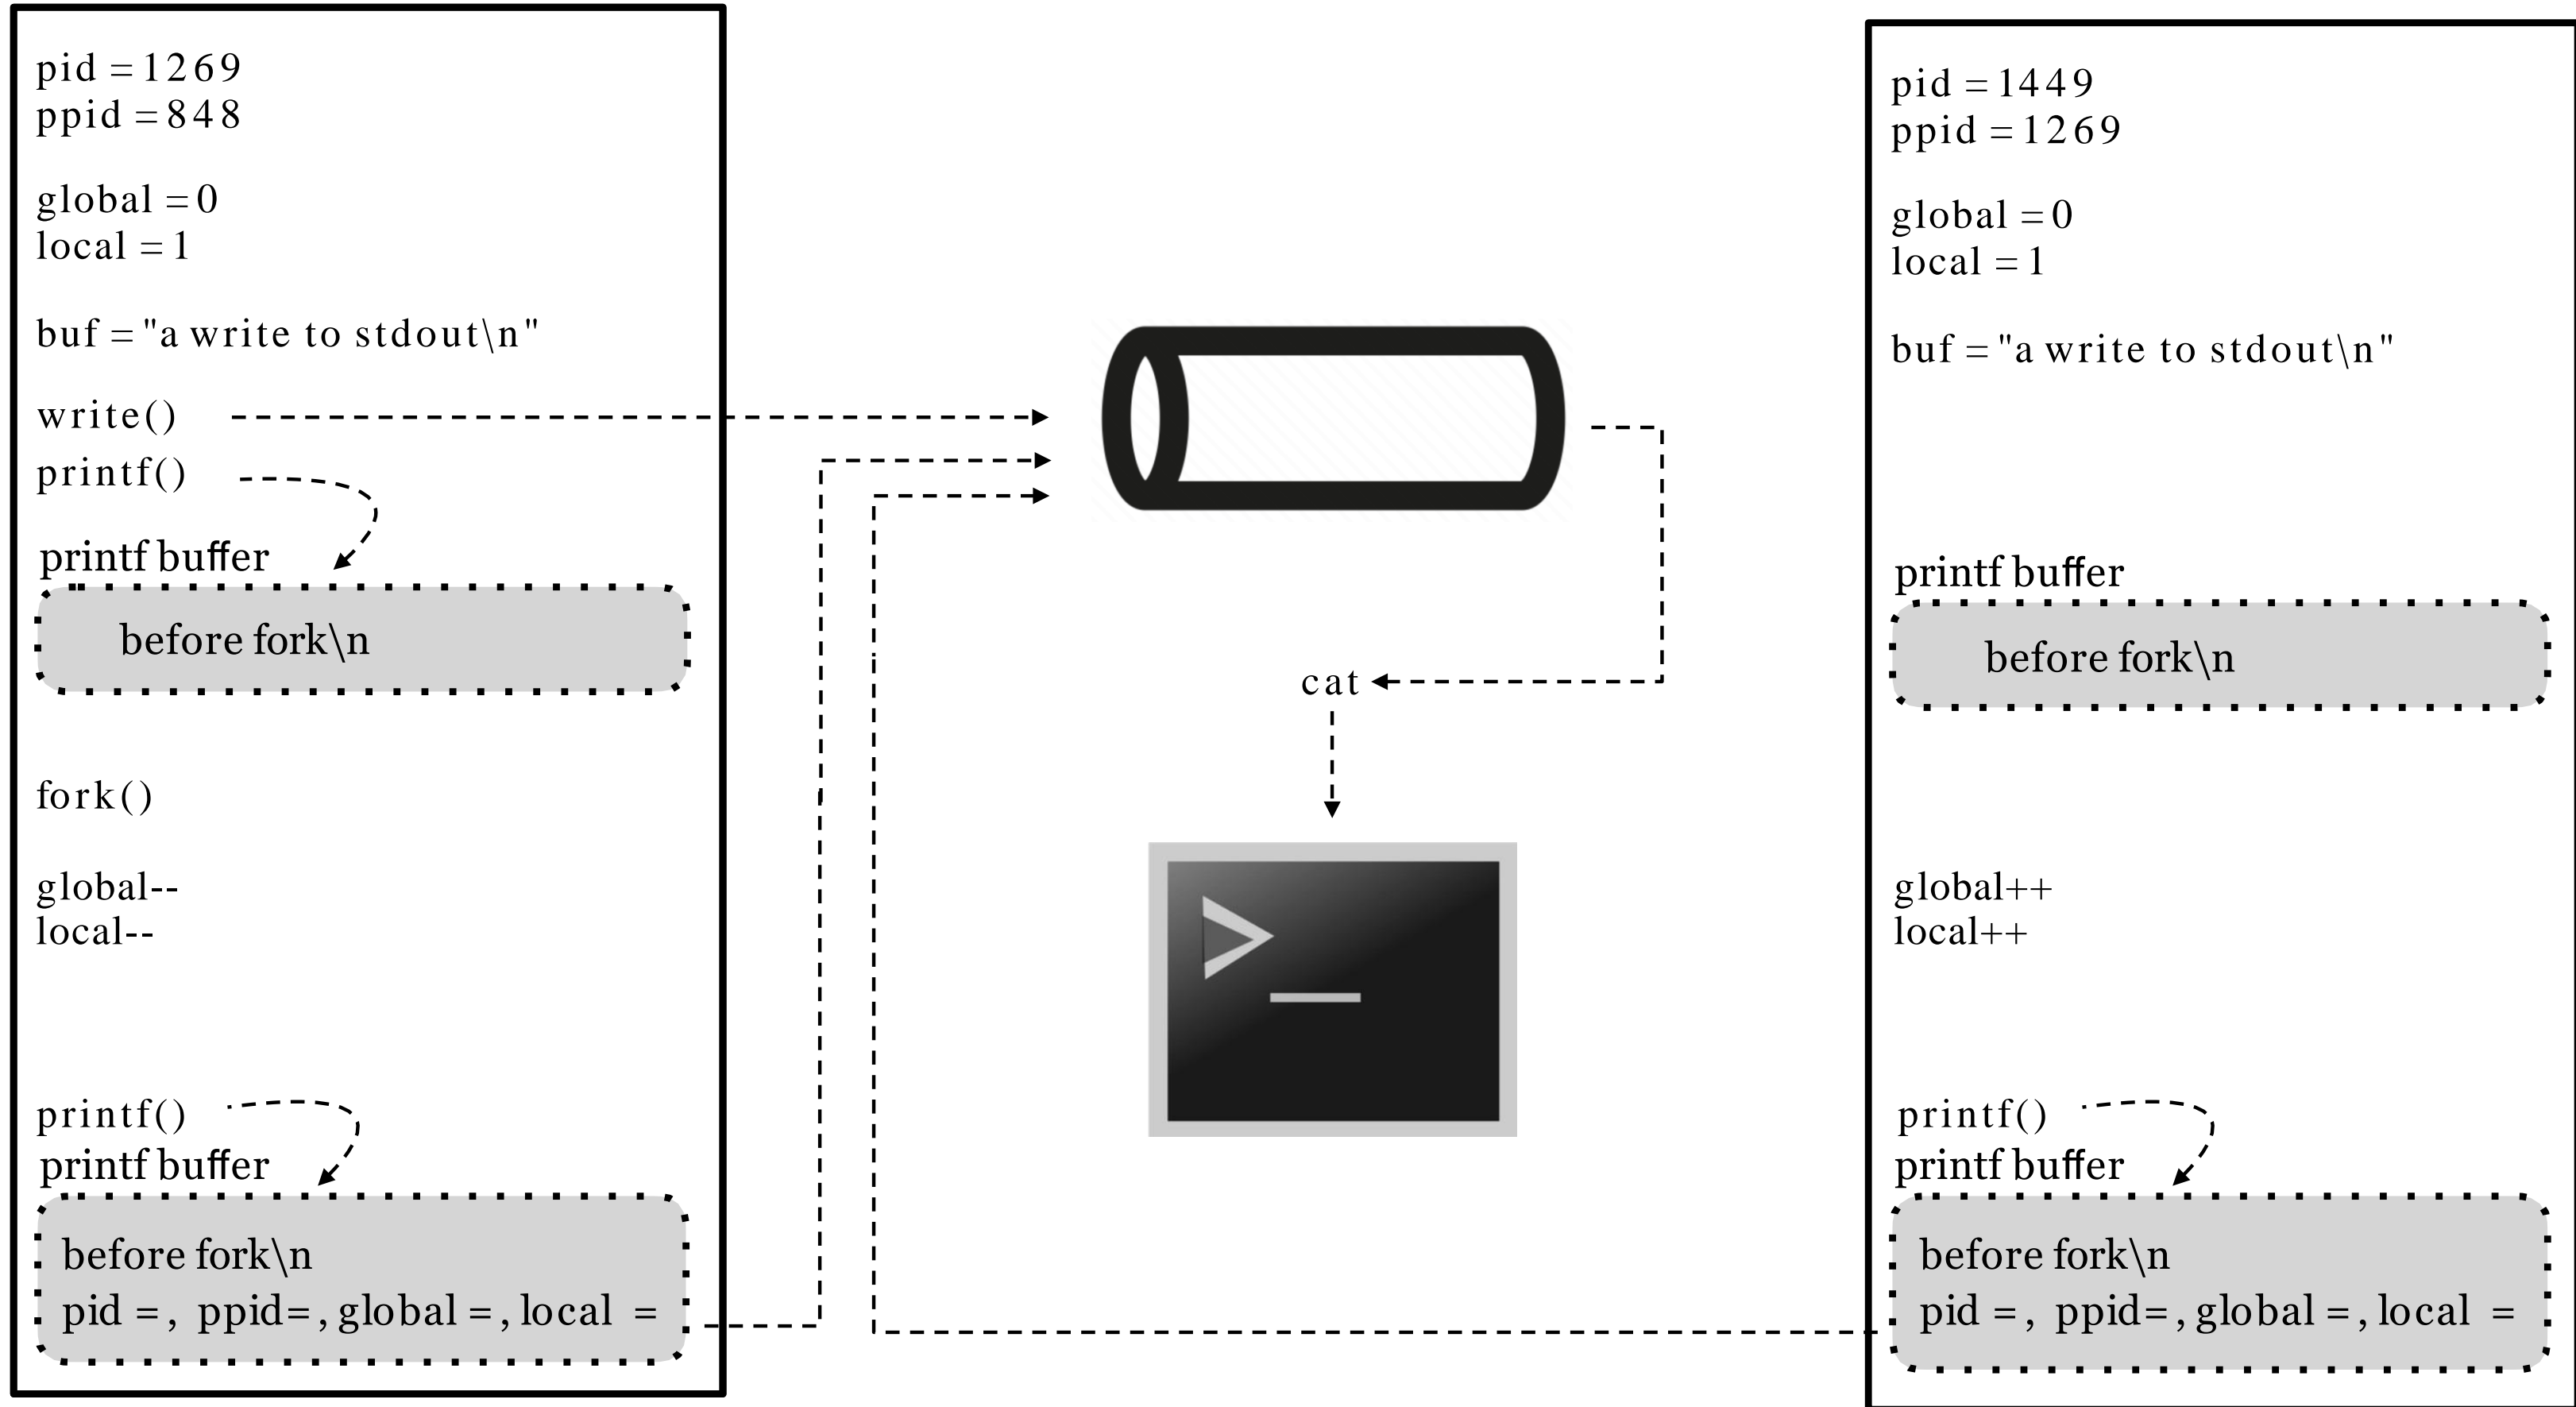
```
[apue$ vi forkflush.c
[apue$ cc -Wall -Werror -Wextra forkflush.c
[apue$ ./a.out
a write to stdout
before fork
pid = 2154, ppid = 1726, global = 1, local = 2
pid = 1726, ppid = 848, global = -1, local = 0
[apue$ echo $$
848
[apue$ ./a.out | cat
a write to stdout
before fork
pid = 1449, ppid = 1269, global = 1, local = 2
before fork
pid = 1269, ppid = 848, global = -1, local = 0
apue$ █
```











## exec(3)

```
#include <unistd.h>

int execl(const char *path, const char *arg, ...);
int execlp(const char *path, const char *arg, ...);
int execlpe(const char *path, const char *arg, ..., char *const envp[]);
int execl_e(const char *path, const char *arg, ..., char *const envp[]);
int execv(const char *path, char *const argv[]);
int execve(const char *path, char *const argv[], char *const envp[]);
int execvp(const char *path, char *const argv[]);
int execvpe(const char *path, char *const argv[], char *const envp[]);
```

Returns: doesn't; -1 on error

The `exec()` family of functions are used to completely replace a running process with a new process image. They all are front-ends for the `execve(2)` system call.

## The `exec(3)` functions

---

- if it has a `v` in its name, `argv`'s are a vector: `const * char argv[]`
- if it has an `/` in its name, `argv`'s are a list: `const char *arg0, .../* (char *) 0 */`
- if it has an `e` in its name, it takes a `char * const envp[]` array of environment variables
- if it has a `p` in its name, it uses the `PATH` environment variable to search for the file
  
- open file descriptors are inherited, unless the `close-on-exec` file flag was set
- ignored signals in the calling process are ignored after `exec`, but caught signals are reset to default
- real `UID/GID` is inherited; effective `UID/GID` is inherited unless the executable was `setuid/setgid`

## wait(2) and waitpid(2)

```
#include <sys/wait.h>

pid_t wait(int *status);
pid_t waitpid(pid_t wpid, int *status, int options);

# include <sys/resource.h>
pid_t wait3(int *status, int options, struct rusage *rusage);
pid_t wait4(pid_t wpid, int *status, int options, struct rusage *rusage);
```

Returns: child PID on success; -1 on error

`wait()` suspends execution of the calling process until status information is available for a terminated child process.

`waitpid()` / `wait4()` allow waiting for a specific process or process group; `wait3()` / `wait4()` allow inspection of resource usage.

## `wait(2)` and `waitpid(2)`

---

Once we get a termination status back in `status`, we'd like to be able to determine how a child died. We do this with the following macros:

- `WIFEXITED(status)` – true if the child terminated normally; use `WEXITSTATUS(status)` to get the exit status
- `WIFSIGNALED(status)` – true if child terminated abnormally (by receiving a signal it didn't catch); use
  - `WTERMSIG(status)` to retrieve the signal number
  - `WCOREDUMP(status)` to see if the child left a core image
- `WIFSTOPPED(status)` – true if the child is currently stopped; use `WSTOPSIG(status)` to determine the signal that caused this

Additionally, `wait(2)` will block until a child terminates; pass `WNOHANG` to `waitpid(2)` / `wait(4)` to return immediately.



# What happens if we don't `wait(2)`?

---





4175 pts/1 Z+ 0:00.00 (a.out)

====

1102 pts/1 Z+ 0:00.00 (a.out)

2668 pts/1 Z+ 0:00.00 (a.out)

2974 pts/1 Z+ 0:00.00 (a.out)

2996 pts/1 S+ 0:00.00 ./a.out

4175 pts/1 Z+ 0:00.00 (a.out)

====

1102 pts/1 Z+ 0:00.00 (a.out)

2668 pts/1 Z+ 0:00.00 (a.out)

2974 pts/1 Z+ 0:00.00 (a.out)

2996 pts/1 S+ 0:00.00 ./a.out

3723 pts/1 Z+ 0:00.00 (a.out)

4175 pts/1 Z+ 0:00.00 (a.out)

I'm going to sleep - try to kill my zombie children, if you like.

[1] Terminated ./a.out

apue\$ ps

PID	TTY	STAT	TIME	COMMAND
701	pts/0	Is	0:00.09	-sh
2262	pts/0	S+	0:00.01	screen -r -d
2134	pts/1	0+	0:00.00	ps
4217	pts/1	Ss	0:00.01	/bin/sh
2388	pts/2	Is+	0:00.01	/bin/sh
2361	pts/3	Ss+	0:00.01	/bin/sh

## Process Control

---

All processes not explicitly instantiated by the kernel were created via `fork(2)`.

`fork(2)` creates a copy of the current process, including file descriptors and output buffers.

To replace the current process with a new process image, use the `exec(3)` family of function.

After creating a new process via `fork(2)`, the parent process can `wait(2)` for the child process to reap its exit status and resource utilization; failure to do so will create a zombie process until the parent is terminated, at which point `init` will reap it.